1 · PRODUCT FOCUS

Catalyst turns one belief about a market event into a few conditional trade plans that the trader reviews and approves.

What it is

A research copilot for retail traders. It maps the belief into competing branches and gives each conclusion one pre-filled plan.

What it is not

Not a chatbot essay and not a signal service. It never places an order: approving a plan is not executing a trade.

Its default answer

WAIT. A plan stays on hold until its entry conditions are observed, because “not yet” is the advice traders skip most.

The model proposes. Code decides. The user approves.

Our customer is the retail trader who acts on events. We give them a plan they can check, not an essay.

WHO

Active retail traders who trade on events and hold positions for weeks. They have a view, but not a plan.

PROBLEM

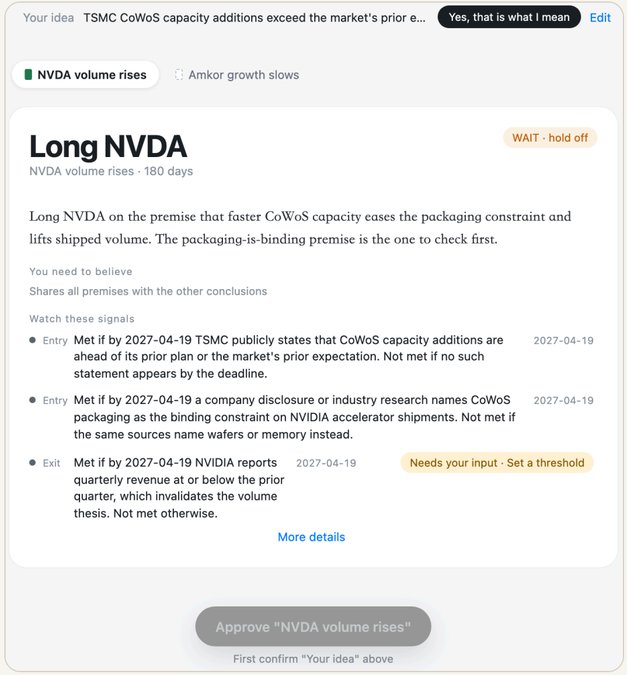
A chatbot answers a headline with 400 balanced words and no decision. The trader still can’t say what must happen first, or what would prove them wrong.

VALUE

Competing branches in one graph, each ending in a plan whose conditions name the data, threshold and deadline. It says WAIT until they are met.



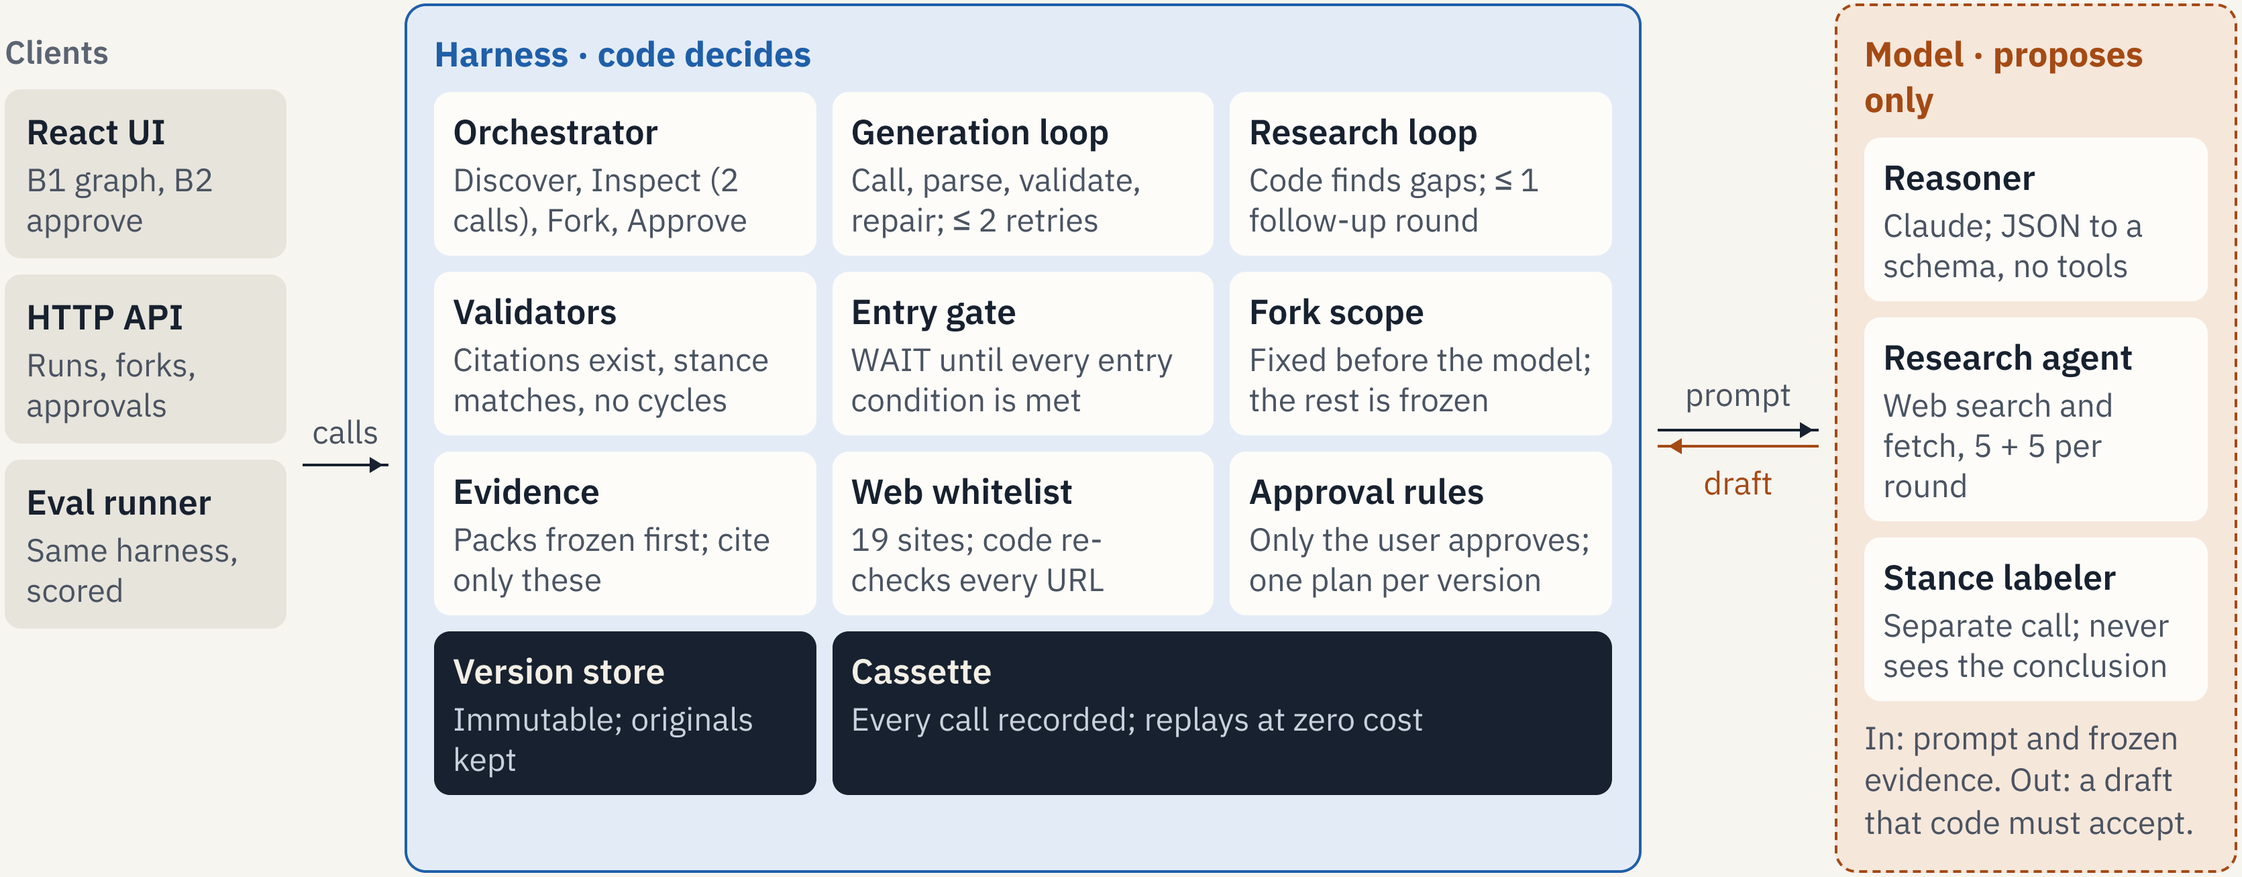
B1 graph. Two competing branches. Solid = backed by sources; dashed = model inference only.



B2 plan. One plan per conclusion, on hold until its entry conditions are observed.

A code harness wraps the agent: the model proposes, code checks.

Left alone, a model cites evidence that isn’t there, approves its own plan, or edits what it shouldn’t.



4 · EVALUATION

78.3 out of 100: the guardrails hold, the content is only fair.

Latest run: Claude Haiku 5.5, prompt p7, offline and online cases, recorded 2026-10-19.

Total score
78.3
out of 100

=

Code checks · half
91.6
out of 100

+

LLM judge · half
6.5
out of 10

Six code checks: share of units passed



What it does not prove

Small samples. The judge saw 19 cases, and research lift rests on a single online case.

An uncalibrated judge. The judge is another model. No human has checked its ratings yet.

An old recording. It predates the fix that draws all candidate directions together; four cases must be re-recorded.

No proof of profit. That needs settled-event backtests and real users.

5 · NEXT STAGE · THE FOUNDER'S THREE HARD PROBLEMS

Next stage: prove the plans are right, that traders use them, and that they can be acted on.

	TODAY	NEXT	HOW WE'LL KNOW
1 · Output quality Are the plans right?	Judge 6.5 / 10. Grounding 51 / 75. Two settled-event backtests (R2, R3), both correct.	Re-record under current prompts. Add human review, which overrides the judge. Run more backtests; add sources, including price data.	Human-reviewed scores and backtest results
2 · UX Will traders use and trust it?	Walked end to end internally. No real users yet.	Test with traders who hold real positions. Add structural graph editing and suggested links, so users can add their own reasoning.	Do users approve plans, edit them, and come back?
3 · Execution Can an approved plan be acted on?	Plans are drafts; approving is not ordering. No market data yet.	Give thresholds structured data sources, mark which conditions code can check, add live market references, export approved plans.	Share of conditions that code can check on its own

Order: quality first, since everything rests on it; user tests in parallel; execution last, since it touches money.